

CIMET QA Gate

Score the sale before it **ships**.

Every sale is checked against its retailer's checklist before it can be submitted. Every verdict points to the exact transcript line, audio timestamp and rule version behind it.

3

outcomes: HELD · QA
· SUBMITTED

21


checks in 2 packs

8

named guardrails

93

self-test assertions



A manual process that can't scale

Today every sale waits for a person to listen to the whole call.

1

Auditor pulls the recording

Downloaded by hand from the dialler

2

Listens end to end

A 30-minute call is 30 minutes of listening

3

Fills an Excel form

No link back to the moment in the call

4

The sale waits

Customers and agents are blocked meanwhile

1,000s

sales a month

~30

retailers, each with its own checklist

30 min

of listening to review one call

0

links from the form back to the audio

"This doesn't scale, and it won't survive a serious retailer audit."

CIMET QA Automation brief

Every sale gets one of three verdicts

Decided before the sale can be submitted, in seconds, with the reason attached.

HELD

A critical check failed

- Routed to a team leader queue
- Card shows said vs expected
- Click the time, hear 20 seconds
- Fix it, or override with a reason

QA

A critical check is uncertain

- Unsure, or confidence below 0.60
- Routed to QA review
- Nothing uncertain auto-passes
- Silence is never treated as a fail


SUBMITTED

Every critical check passed

- Cleared to submit to the CRM
- No human touch needed
- 5% still sampled for human audit
- Measures the model, not just trusts it

From hang-up to verdict, with zero manual handling

Built in the order the brief asked for: recording ingestion first.



- **Asynchronous:** the webhook answers 202 at once; the dialler never waits on transcription.
- **Nothing lost:** the recording is saved before speech-to-text runs, so a failed transcription never loses the call.
- **Honest:** with no speech-to-text key the webhook refuses (503) - it never passes a canned transcript off as the call.

Recording ingestion: the dialler webhook

POST /api/dialler/recording — recording in, keyed on Lead ID, scored automatically.

```
# raw audio body - what a dialler integration sends
POST /api/dialler/recording?lead_id=3613802
Content-Type: audio/wav      <binary>

# or a URL the server fetches
{"lead_id": "3613802", "recording_url": "https://..."}

# or base64 in JSON
{"lead_id": "...", "audio_base64": "...", "filename":
"call.wav"}

-> 202 {job_id, poll: /api/jobs/<job_id>}
```

- **Speaker separation + timestamps** from ElevenLabs Scribe or Deepgram, chosen by which key is set.
- **Agent found by what they say** - compliance phrases only an agent uses - not by who spoke first. Method and margin are recorded.
- **Card numbers masked before disk** - 13+ digits, whether ASR wrote “4111...” or “four one one one...”.
- **Recording kept + checksummed** - SHA-256 on the transcript, served back with seek support for click-to-play.
- **Optional shared secret** - set CIMET_DIALLER_TOKEN and the dialler must send X-Dialler-Token.
- **Proven on real audio** - calls rendered with ElevenLabs and transcribed by Scribe through this exact webhook.



Every transition is appended to runs/dialler_log.jsonl.

Script, facts, behaviour — each scored its own way

The part the brief said matters most.



Verbatim / script

Transcript vs approved script

Token-level match over a sliding window of turns, tightened to the moment it was said. Handles ASR run-together words and alternative wordings.

LLM No LLM. No hallucination surface.

BLOCKS SALE **Yes, if critical**

EXAMPLES Recording disclaimer + consent, account holder, DMO read verbatim



Factual match

Transcript vs CRM fields vs plan

The extractor reports what was SAID (quote + time). Deterministic code compares it to the CRM or plan value.

LLM Extraction only. Comparison is code.

BLOCKS SALE **Yes**

EXAMPLES Rate quoted, email, NMI, fuel type, promo & ongoing price, speed, TMC, fees, contract term



Behaviour

Transcript only

Arithmetic over turn timings and wording: longest silence, talk-over, cut-off turns, customer confusion.


LLM No LLM.

BLOCKS SALE **Never**

EXAMPLES Dead air, interruptions, talk-over, customer confusion

The LLM never sees the answer

It reports what was said. Code decides whether that matches.



Structural

`build_prompt()` takes the check and the transcript. CRM and plan data are not in scope.

Runtime leak guard

A prompt containing a reference value not spoken on the call is refused before sending.

Grounded quotes (G3)


A quote that isn't in the transcript is discarded. Invented evidence can't reach a verdict.

Audited

Model, effort, prompt SHA-256 and extractor are stamped on every result.

When a sale actually goes through

Rules evaluated in order; the first match wins.



Built in, not configured

- **Type C removed first:** a coaching note can't change the outcome, even if someone flips it.
- **Real queue:** HELD and QA append a routing record to runs/queue.jsonl.
- **Re-runs live:** every override re-evaluates the gate and logs before → after.
- **Floor is a setting:** CIMET_CONFIDENCE_FLOOR, 0.60 by default.

The brief: “Any critical fail is held, low-confidence scores route to a human, nothing uncertain auto-passes.”

Eight guardrails that stop false criticals

Each one is named on the result it changed — in the console and the audit file.

G1**Audio-quality abstain**

A script span that would fail, but sits in [inaudible]/[crosstalk] or ASR < 0.75, becomes unsure. Unheard is not unsaid.

G2**Abstention honoured**

When the extractor can't establish a value, the check is unsure — never a fail.

G3**Ungrounded quote**

A quote that isn't in the transcript is discarded. Hallucinated evidence can't reach a verdict.

G4**Not comparable**

A value that won't parse, or no reference on the lead, is unsure rather than guessed.

G5**Low-confidence critical**

A critical fail below the 0.60 floor becomes unsure: QA, not a hold.

G6**Timestamp corrected**

If the model's time disagrees with its quote's turn, the transcript wins and it's logged.

G7**Possible mishear**

Emails audio can't settle (rahman/raman, bigpond heard as bizpoint) go to a human. CRM typos like gmial still fail.

G8**Screen vs spoken**

Agent's words vs the order screen: a FAIL at 0.7× confidence; the critical floor decides FAIL or QA.

Built for how a team leader actually works

See the verdict → see what's wrong → hear it → fix it or override.

CIMET QA Gate

3613790 Janet Smith

ASR elevenlabs Type B claude-opus-5

Accuracy

Re-run

HELD

2 critical checks failed

Peak usage rate stated matches the sold plan - Customer email captured correctly. Held and routed to team leader review.

5 Passed

2 Failed

0 Unsure

65% Score

Checks

RI-ENERGY-SALES-QA RI-2026.09 - 9 checks

Needs attention 2

Peak usage rate stated matches the sold plan

Failed - Rates and charges

SAID 28.6

EXPECTED 31.9

spoken 28.6c/kWh is 3.3c below the plan rate of 31.9c/kWh

CRITICAL TYPE B conf 0.94

Override

Customer email captured correctly

Failed - Rates and charges

SAID j.smith@gmail.com

EXPECTED j.smith@gmial.com

spoken j.smith@gmail.com vs CRM j.smith@gmial.com - 2 character differences; the CRM domain gmial.com is not a real provider and is 2 keystrokes from gmail.com: a keying error in the CRM

CRITICAL TYPE B conf 0.92

Override

Passed 5

Coaching notes 2

How the three check types are scored

Timeline

24:57 · 96 turns · no recording · whisper-large-v3

00:00 06:14 12:29 18:43 24:57

agent customer fail unsure pass dead air

Transcript CRM & plan Audit Overrides

1 card-length number masked before anything was stored or scored.

00:00 Daniel R. Good afternoon, you're speaking with Daniel from Retailer One. Before we go any further I need to let you know that this call is being recorded for quality assurance and training purposes. Do I have your consent to continue on a recorded line?

00:14 Janet S. Yeah, that's fine, go ahead.

00:18 Daniel R. Thank you. Can I just confirm I'm speaking with the account holder, Janet Smith, and that you're authorised to make changes to this account?

00:32 Janet S. Yes, that's me. Janet Smith. I'm the account holder, the account's in my name.

00:39 Daniel R. Perfect, thank you Janet. And can I confirm the supply address is 12 Ruskin Street, Elwood, Victoria, three one eight four?

00:52 Janet S. That's the one.

00:56 Daniel R. Great. So this is the follow-up on the Saver Flex plan we spoke about last week. Before I take you through any numbers there's a short statement I'm required to read to you in full, so bear with me for a moment.

01:17 Janet S. No worries.

- 1

Verdict

HELD / QA / SUBMITTED, which checks and where it was routed, plus pass / fail / unsure counts.
- 2

Needs attention

Failed and unsure checks as cards: said vs expected, the reason, Play 20s, Override.
- 3

Call timeline

Agent and customer lanes, a pin per check, dead air hatched. Click to jump or play.
- 4

Tabs

Transcript, CRM & plan (read-only), Audit, Overrides.
- 5

Lead picker + engines

Each lead's latest verdict as a coloured dot; ASR and extractor status.

Every lead's verdict at a glance

One button, one grouped list — no hunting through tabs.

The screenshot shows the 'CIMET QA Gate' interface for 'Retailer 1 - checklist R1-2026.09'. A dropdown menu is open, showing a list of leads. The leads are categorized into 'Brief leads' and 'Synthetic test calls'. Each lead entry includes a status dot, a lead ID, a name, and a brief description. The status dots are color-coded: red for 'HELD', yellow for 'QA', and green for 'SUBMITTED'. The dropdown also shows a 'Checks' section with 'Needs attention' (2) and a list of failed checks, including 'Peak usage rate stated match' and 'Customer email captured correctly'.

CIMET QA Gate
Retailer 1 - checklist R1-2026.09

HELD 2 critical checks
Peak usage rate stated match

Checks
Needs attention 2

✗ **Peak usage rate stated match**
Failed · Rates and charges

SAID
28.6
spoken 28.6c/kWh is 3.3c below the

CRITICAL TYPE B conf 0.94

✗ **Customer email captured correctly**
Failed · Rates and charges

SAID

Brief leads

- **3613790** Janet Smith
Worked example from the brief: rate quoted 28.6 vs plan 31.9 at 14:02, email mismatch
- **3613791** David Nguyen
Clean compliant call. Expected gate: SUBMITTED (subject to 5% sample audit).
- **3613792** Mario Rossi **NOISY**
Poor audio, heavy crosstalk. Silence is NOT evidence of non-compliance. Expected ...
- **3613793** Redacted customer **NBN**
Official CIMET artefact: redacted internet (NBN) sales call, scored against the NBN-...

Synthetic test calls

- **3613801** Tom Becker
Clean compliant call
- **3613802** Aisha Rahman **ASR**
Wrong rate quoted (27.9 vs plan 31.9)
- **3613803** Liam O'Connor
Email keyed wrong in CRM
- **3613804** Grace Palmer

The dot is the verdict

- **HELD** a critical check failed
- **QA** a critical check needs a human
- **SUBMITTED** cleared

Tags



NBN-QA checklist · scored from a real dialler recording · poor audio

- **Grouped:** the four brief leads, then the 12 synthetic test calls, each with a one-line note.
- **Always coloured:** never-scored leads are scored once in the background at startup (saved, never repeated).
- **Keyboard:** arrow keys move, Escape closes; any press outside closes it.

Said vs expected, in one glance

The worked example from the brief, exactly as the console shows it.

Checks

R1-ENERGY-SALES-QA R1-2026.09 · 9 checks

Needs attention 2

✗ Peak usage rate stated matches the sold plan

Failed · Rates and charges

14:02

SAID

28.6

≠

EXPECTED

31.9

spoken 28.6c/kWh is 3.3c below the plan rate of 31.9c/kWh

CRITICAL

TYPE B

conf 0.94

Override

✗ Customer email captured correctly

Failed · Rates and charges

22:10

SAID

j.smith@gmail.com

≠

EXPECTED

j.smith@gmial.com

spoken j.smith@gmail.com vs CRM j.smith@gmial.com - 2 character differences; the CRM domain gmial.com is not a real provider and is 2 keystrokes from gmail.com: a keying error in the CRM

CRITICAL

TYPE B

conf 0.92

Override

Passed 5

Coaching notes 2

► How the three check types are scored

Checks

R1-ENERGY-SALES-QA R1-2026.09 · 9 checks

CRITICAL

TYPE B

conf 0.94

Override

✗ Customer email captured correctly

Failed · Rates and charges

22:10

SAID

j.smith@gmail.com

≠

EXPECTED

j.smith@gmial.com

spoken j.smith@gmail.com vs CRM j.smith@gmial.com - 2 character differences; the CRM domain gmial.com is not a real provider and is 2 keystrokes from gmail.com: a keying error in the CRM

CRITICAL

TYPE B

conf 0.92

Override

Passed 5

Recording disclaimer given and consented

00:00

Account holder identity and authority confirmed

00:18

DMO / reference price statement read verbatim

01:21

NMI read back and confirmed

22:35

Fuel type confirmed

23:29

Coaching notes 2

Dead air on the call

26s of dead air at 18:30

18:30

Talk-over and interruptions

2 talk-over event(s), first at 13:44 - within the 3 coaching threshold

13:44

► How the three check types are scored

SAID

28.6

≠

EXPECTED

31.9

c/kWh at 14:02 — 3.3c below the plan sold

- **Every card has:** the check, section, time, said vs expected, one-line reason, CRITICAL / TYPE / confidence, Play 20s and Override.
- **Missing script wording** is listed phrase by phrase with how much of it was found.
- **Sorted for action:** critical first, fails before unsure, then call order.
- **Passed + coaching notes** fold away (right) — one click from their own evidence.

Every verdict shows its working

Nothing a team leader is asked to trust is hidden.

Checks

R1-ENERGY-SALES-QA R1-2026.09 · 9 checks

daily supply charge is 98.45 cents a day.

spoken value	28.6 cents per kilowatt hour (including GST) (extraction confidence 0.97)
reference	Plan peak rate = 31.9c from plan
comparison	mismatch - spoken 28.6c/kWh is 3.3c below the plan rate of 31.9c/kWh
extractor	llm:claude-opus-5 (effort medium)
leak guard	the extractor never saw Plan peak rate
grounding	quote found in the transcript (100% match)
prompt sha256	1e14f367b571b212e688da7d...
check version	1.4.0 (from 2026-09-01)
basis	NECF Rule 46 - price disclosure before consent

Customer email captured correctly

Failed · Rates and charges

22:10

SAID

j.smith@gmail.com

≠

EXPECTED

j.smith@gmial.com

spoken j.smith@gmail.com vs CRM j.smith@gmial.com - 2 character differences; the CRM domain gmial.com is not a real provider and is 2 keystrokes from gmail.com: a keying error in the CRM

CRITICAL

TYPE B

conf 0.92

Override

Passed 5

Coaching notes 2

How the three check types are scored

Quote

verbatim line, speaker, turn and time

Spoken value

what the extractor heard + its confidence

Reference

the CRM / plan value and its source

Comparison

match or mismatch, and exactly why

Extractor

Claude model + effort, or the fallback — and why

Leak guard

confirms the model never saw the reference

Grounding

quote found in the transcript (and how closely)

Prompt SHA-256

the exact prompt, reproducible

Check version

e.g. 1.4.0, in force from 2026-09-01

Basis

e.g. NECF Rule 46 — price disclosure before consent

Click the timestamp, hear twenty seconds

“Nobody listened to thirty minutes to find that.” — the brief's worked example



Two lanes

Agent on top, customer below — who spoke when, at a glance.



A pin per check

Red fail, amber unsure, green pass.
Click: open the check + play it.



Dead air hatched

Long silences shown where they happened; coaching only.



Playhead + highlight

Follows the recording; the line being spoken lights up.

Scribe transcript → Claude extraction → held in seconds

Lead 3613802: recording pushed through the webhook, transcribed by ElevenLabs Scribe.

CIMET QA Gate

Retailer 1 · checklist R1-2026.09

3613802

Aisha Rahman

ASR

ASR elevenlabs

Type B claude-opus-5

Accuracy

Re-run

HELD

1 critical check failed

Peak usage rate stated matches the sold plan. Held and routed to team leader review.

5

Passed

1

Failed

1

Unsure

72.5%

Score

Checks

R1-ENERGY-SALES-QA R1-2026.09 · 9 checks

✖

Peak usage rate stated matches the sold plan

01:38

Failed · Rates and charges

SAID

27.9

EXPECTED

31.9

≠

spoken 27.9c/kWh is 4c below the plan rate of 31.9c/kWh

CRITICAL

TYPE B

conf 0.97

▶ Play 20s

Override

01:38 · agent · turn 13

On the Saver Flex plan, your usage rate is twenty-seven point nine cents per kilowatt hour including GST, and the daily supply charge is ninety-eight point four five cents a day.

spoken value

27.9 cents per kWh including GST (extraction confidence 0.97)

reference

Plan peak rate = 31.9c from plan

comparison

mismatch - spoken 27.9c/kWh is 4c below the plan rate of 31.9c/kWh

extractor

llm:claude-opus-5 (effort medium)

leak guard

the extractor never saw Plan peak rate

grounding

quote found in the transcript (100% match)

prompt sha256

46a544f2308406edf49a4bdb...

check version

1.4.0 (from 2026-09-01)

basis

NECF Rule 46 - price disclosure before consent

?

Customer email captured correctly

02:13

Unsure · needs a human · Rates and charges

▶ How the three check types are scored

Call timeline

02:53 · 26 turns · from dialler recording · elevenlabs:scribe_v1

00:00

00:43

01:27

02:10

02:53

▶ 1:42 / 2:54

3613802.wav · sha256 1776da53c2...

Transcript

CRM & plan

Audit

Overrides 0

Okay, got it.

01:23 Agent (speaker_0) A couple of quick questions first. Does anyone at the property rely on life support equipment, and do you hold a concession card?

01:33 Customer (speaker_1) No to both.

01:35 Agent (speaker_0) Let me pull the pricing up for you.

01:38 Agent (speaker_0) On the Saver Flex plan, your usage rate is twenty-seven point nine cents per kilowatt hour including GST, and the daily supply charge is ninety-eight point four five cents a day.

01:50 Customer (speaker_1) That sounds reasonable.

01:53 Agent (speaker_0) The rate isn't locked. We'd give you written notice before any change, and you could leave then with no exit fee. Are you happy for me to go ahead?

02:03 Customer (speaker_1) Yes, let's do it.

02:06 Agent (speaker_0)

- **Source:** “from dialler recording · elevenlabs:scribe_v1”, with the file's SHA-256.
- **Rate:** said 27.9, plan 31.9 — critical fail at 01:38, confidence 0.97.
- **Email:** unsure (G7) — Scribe wrote “raman” for “rahman”, which audio can't settle.
- **Verdict:** HELD, routed to team leader review.

CIMET QA Gate

15

The record and the audit trail, one click away

CRM & plan are read-only — the app never writes to the CRM.

CRM record

read-only

account holder name	Janet Smith
email	j.smith@gmial.com
phone	0412 884 210
nmi	61020345678
fuel type	electricity
site address	12 Ruskin Street, Elwood VIC 3184
concession card	no
life support	no

Plan sold

R1-SAVER-FLEX

plan	Retailer 1 Saver Flex
peak rate	31.9 c/kWh
daily supply	98.45 c/day
feed-in	4.5 c/kWh
vs DMO	-4%
term	no lock-in

CRM & plan: read-only; the failed field is flagged — the CRM holds j.smith@gmial.com.

Redacted leads stay redacted

On the NBN artefact the CRM record shows the placeholder tags exactly as supplied - nothing is re-identified or filled in.

[CUSTOMER_FULL_NAME]

[EMAIL]

[PHONE]

[DOB]

[SERVICE_ADDRESS]

Transcript

The whole call; the line being spoken is highlighted; any timestamp plays from there.

CRM & plan

Read-only. A field that failed is flagged in red. Redacted tags stay as tags.

Audit

Run ID, checklist version, extractor, floor, score, coverage, checksum, pipeline trace, audit JSON.

Overrides

The append-only log for this run: who, why, machine verdict, gate before → after.

Lead 3613790: the brief's example, reproduced

Retailer 1, 25-minute call. Every verdict matches the brief.

Check	The brief expects	CIMET QA Gate shows
Recording disclaimer + consent	PASS	PASS · 00:00
Account holder confirmed	PASS	PASS · 00:18
DMO read verbatim	PASS	PASS · 01:21
Rates and charges	FAIL · 28.6 vs 31.9 · 14:02	FAIL · 14:02 · said 28.6 ≠ 31.9
Email captured	FAIL · gmail vs gmial · 22:10	FAIL · 22:10 · gmial.com is a CRM keying error
Dead air (coaching)	NOTE · 47s at 18:30	NOTE · dead air at 18:30 · never blocks

HELD

2 critical checks failed. Routed to team leader review.

The TL clicks 14:02, hears twenty seconds, corrects the email and re-confirms the rate — or the sale is cancelled.

Our transcript's opening checks fall at different times than the brief's illustration; the verdicts are identical.

Silence is never evidence

Lead 3613792: crosstalk, [inaudible], a dropped line → QA, not HELD.

Checks

R1-ENERGY-SALES-QA R1-2026.09 · 9 checks

of compliance

The Default [inaudible] Offer is the [crosstalk] set by the Australian Energy [inaudible], and the estimated annual cost of this plan is- Sorry mate, you've gone all choppy. Say again? [inaudible] per cent less than the- for an average household in your- sorry, the line's terrible. [crosstalk]

G1 audio-quality abstain: coverage 44% would normally fail, but 5 audio marker(s) (crosstalk, inaudible) in the span and ASR confidence as low as 0.38 mean the wording may have been spoken and not captured. Downgraded to unsure for human review.

✗ "Default Market Offer"67%

✗ "reference price"0%

✗ "Australian Energy Regulator"67%

✗ "this is an estimate only"60%

✗ "your actual cost will depend on your usage"25%

methoddeterministic script-span match, no LLM

coverage44% (pass at 85%, unsure at 55%)

audio markerscrosstalk, inaudible

asr confidencemin 0.38 / mean 0.51

check version2.0.0 (from 2026-07-01)

basisAER Retail Pricing Information Guidelines - Default Market Offer comparison

? Customer email captured correctly06:28

Unsure - needs a human · Rates and charges

SAID

?

EXPECTED

► How the three check types are scored

44%

of the DMO wording found — on its own, a fail

G1

5 audio markers + ASR 0.38 in that span → unsure, not fail

G2

rate and email obstructed → extractor abstains, never guesses

QA

a human decides; audible checks still pass

The judging criterion: “Handles messy real calls — crosstalk, mishears, silence — without manufacturing false criticals.”

A mishear is not a keying error

ElevenLabs Scribe heard “bigpond” as “bizpoint”. Before G7, that held a clean sale.

Checks

R1-ENERGY-SALES-QA R1-2026.09 · 9 checks

?

Customer email captured correctly

02:20

Unsure - needs a human · Rates and charges

SAID

mario.bianchi@bizpoint.com

EXPECTED

mario.bianchi@bigpond.com

the CRM holds a real provider (bigpond.com) but the transcript has bizpoint.com, which is not one - most likely misheard on the call.

CRITICAL

TYPE B

conf 0.47

Play 20s

Override

02:20 · customer · turn 19

It's mario.bianchi@bizpoint.com.

G7 possible mishear: the CRM holds a real provider (bigpond.com) but the transcript has bizpoint.com, which is not one - most likely misheard on the call. Reported as unsure so a person listens to the clip, rather than holding the sale on how ASR spelled it.

spoken value

mario.bianchi@bizpoint.com (extraction confidence 0.86)

reference

CRM email = mario.bianchi@bigpond.com from crm

comparison

mismatch - spoken mario.bianchi@bizpoint.com vs CRM mario.bianchi@bigpond.com - 3 character differences; the CRM holds a real provider (bigpond.com) but the transcript has bizpoint.com, which is not one - most likely misheard on the call

HEARD mario.bianchi@bizpoint.com

CRM mario.bianchi@bigpond.com

UNSURE → human

Real provider in CRM, non-existent one heard

HEARD aisha.raman@gmail.com

CRM aisha.rahman@gmail.com

UNSURE → human

Spellings sound identical

HEARD j.smith@gmail.com

CRM j.smith@gmial.com

FAIL → held

CRM domain isn't a real provider: keying error

Also from real ASR: US spellings (“authorized”) now match Australian scripts (“authorised”).

Checklist packs — and the official NBN artefact

Lead 3613793: the redacted internet sales call, scored against a 12-check NBN pack.

CIMET QA Gate
Provider A (internet) · checklist NBN-2026-09

3613793 Redacted customer NBN

ASR: elevenlabs · Type B: claude-opus-5 · Accuracy · Re-run

QA 3 critical checks need a human
Total minimum cost disclosed correctly · New development fee: spoken guarantee vs order screen · Contract term stated consistently and matches the plan. Routed to QA review - nothing uncertain auto-passes.

7 Passed 0 Failed 3 Unsure 85% Score

Checks PA-NBN-SALES-QA NBN-2026-09 12 checks
Needs attention 3

Contract term stated consistently and matches the plan 02:06
Unsure · needs a human · Costs and fees
SAID: month to month EXPECTED: not supplied
the lead carries no contract term to compare against.
CRITICAL TYPE B conf: 0.46 Override

Total minimum cost disclosed correctly 06:07
Unsure · needs a human · Costs and fees
SAID: \$42.90 EXPECTED: \$317.00
the agent's TMC \$42.90 is the order TMC \$317.00 less the \$275.00 new development fee, which the lead marks not applicable - but the order screen still totals \$317.00 and the agent told the customer to disregard it.
CRITICAL TYPE B conf: 0.55 Override

New development fee: spoken guarantee vs order screen 11:10
Unsure · needs a human · Costs and fees
SAID: \$275.00; not charged EXPECTED: \$275.00; not
How the three check types are scored

Call timeline NBN-120 turns · no recording · source transcript (redacted PDF, provider unknown)
00:00 04:00 08:17 12:25 16:33
agent customer fail unsure pass dead air

Transcript CRM & plan Audit Overrides @

00:00 Speaker 1 Hello, [CUSTOMER_NAME] speaking.
00:02 Speaker 2 Yes, Hi, [CUSTOMER_NAME]. Good day. This is [AGENT_NAME] from Internet's comparison. How are you?
00:08 Speaker 1 Good. Thanks. How are you?
00:10 Speaker 2 Yeah, I'm good. Thank you. And we noticed that you're looking for better Internet plans, and we're calling to assist you with this. Yeah. K. And as I check it, your address is [SERVICE_ADDRESS]. Correct. [UNCLEAR_NAME]. Yep. [SERVICE_ADDRESS]. K. By the way, please be advised that this call will be recorded for quality assurance and training purposes. K? Now let me double check the address. Bear with me. There you go. As I check it, yeah. The address is already an NBN ready, fiber to the premises technology. Okay? Now just to ask, [CUSTOMER_NAME], who's your current Internet service provider? Do you have one? PRIMUS. You are currently with PRIMUS. How much are you paying?
00:52 Speaker 1 They reduced it today to sixty five. That's all right here.
00:56 Speaker 2 Sixty five dollars for how many MBPS? Do you know the speed? How do you care? Twenty five. Twenty five Mbps? Yeah. Okay. And how many people are using the Internet? Just myself. Are you used to Internet? Are you do gaming, video streaming?
01:11 Speaker 1 Work from home? Yeah, I do. No, I don't work from home. I just watch, you know, Netflix and that sort of stuff, and I've seen my TV still at, like, channel seven channel

- **Packs:** Retailer 1 energy (9 checks) is the default; NBN-QA (12) applies to leads that name it.
- **Transcript:** copied row for row from the redacted PDF; all 15 placeholder tags kept verbatim.
- **No regressions:** 0 verdict changes on 3613790–92 after adding the pack.
- **Honest timing:** the PDF has no audio, so timestamps are estimated from turn order — and labelled as such.

Type A · script & rules

Recording disclaimer at the open
Identity verified (name, email, DOB)
Card kept off the recording (mute, no PAN)

Type B · facts vs plan

Promo price + months
Ongoing price
Typical evening speed
Modem model + cost
Total minimum cost
Development fee vs order screen
Contract term

Type C · coaching

Talk-over / cut-off turns
Customer confusion

\$42.90 said. \$317 on the screen.

G8 scores the conflict; the confidence floor decides it goes to a human.

Checks

PA-NBN-SALES-QA NBN-2026.09 · 12 checks

your name, dismissed or missus, or do you have any title?

G8 screen vs spoken conflict: the agent's TMC \$42.90 is the order TMC \$317.00 less the \$275.00 new development fee, which the lead marks not applicable - but the order screen still totals \$317.00 and the agent told the customer to disregard it. The call alone cannot show which the provider will bill, so this is a FAIL at 70% of normal confidence (0.81 -> 0.57); the critical floor decides whether it holds the sale or routes it to QA.

G5 low-confidence critical fail: confidence 0.57 is below the critical floor 0.60, so the fail was downgraded to unsure rather than holding the sale

spoken value	42.90 (extraction confidence 0.9)
reference	Lead TMC and conditional fee = tmc_cents = \$317.00; new_development_fee_cents = \$275.00; fee_applicable = no; promo_price_cents = \$42.90 from lead
comparison	mismatch - agent told the customer the TMC is \$42.90; the order's TMC is \$317.00
extractor	llm:claude-opus-5 (effort medium)
leak guard	the extractor never saw Lead TMC and conditional fee
grounding	quote found in the transcript (100% match)
prompt sha256	216c78c52abfc6cb0b08f4f7...
check version	1.0.0 (from 2026-09-01)
basis	TCP Code - total minimum cost must be disclosed before the sale

? New development fee: spoken guarantee vs order

11:10

► How the three check types are scored

\$42.90

the agent's TMC at 06:07 — “forty two dollars and ninety only”

\$317

read off the order screen 3x — “don't worry about that”

\$275

development fee — “I really guarantee you it will not be charged”

$\$317 - \$275 = \$42 \rightarrow$ the agent quoted the order total without a fee the lead marks not applicable, while the submitted order still includes it.

G8

screen vs spoken: FAIL at 0.7× (0.81 → 0.57)

G5

0.57 < 0.60 floor → unsure

QA

a human confirms the fee is removed before submitting

Contract term: “month to month” and “one to one contract”, and no term on the lead → unsure (none invented).

Humans decide — and every decision is kept

An override changes the gate, never the CRM.

Override check

Peak usage rate stated matches the sold plan (rates_and_charges v1.4.0)
machine verdict: FAIL at 0.94 confidence, 14:02

Set status to

PASSFAIL

Reviewer

TL Sam Okafor

Reason **required**

Listened at 14:02 - agent corrected to 31.9c before consent. Customer confirmed.

At least 10 characters. The machine verdict is kept alongside your decision.

Overrides change the gate only. The CRM record is never written to.

CancelSave override

● A written reason

At least 10 characters — enforced in the UI and again at the API.

● Machine verdict kept

Original status and confidence stay next to the human decision.

● Gate re-runs live

HELD → SUBMITTED only when the last blocking check clears.

● Append-only log

runs/override_log.jsonl: who, why, when, gate before → after.

● CRM never written

crm_written=false on every entry. Fixing the record is the TL's job.

● Coaching can't block

Type C notes can't be overridden into a blocking state.

Every verdict can be rebuilt from one file

runs/<run_id>.json — the Download audit JSON button.

```
# every result
check_id, check_version, effective_from
status, confidence, blocks_sale
quote, start_sec, turn_idx, speaker
guardrails: [G1 ... G8 that fired]
evidence: extractor, model, prompt_sha256,
          reference + source, comparison

# every run
checklist version in force
transcript_sha256 (+ recording sha256)
crm_snapshot, plan_snapshot
gate decision + reasons + history
override_log, pipeline_trace, config
```

The screenshot shows the 'Audit' tab of the CIMET QA Gate interface. It displays the following information:

- Run ID:** RUN-3613790-20260919T085112-5d6b
- Checklist:** R1-ENERGY-SALES-QA R1-2026.09
- LLM:** claude-opus-5 (Type B extractor)
- Confidence:** 0.60 (critical confidence floor)
- Weighted Score:** 65%
- Evidence Coverage:** 100%
- Queue:** TEAM_LEADER_REVIEW
- Transcript SHA256:** e5143bfe12ad0d...

A table below the audit details shows the following steps and durations:

Step	Description	Duration
ingest	96 turns, 24:57 of audio, 1 redaction(s) applied	0ms
extract	4 Type B extraction(s) via llm:claude-opus-5	3235ms
score	9 checks scored against checklist R1-2026.09	153ms
gate	HELD -> TEAM_LEADER_REVIEW	0ms

A 'Download audit JSON' button is located at the bottom of the audit details section.

Rule versions

- **Stamped today:** every result carries its check version and the date it took effect.
- **Next step:** resolve each check to the version that was live on the call date, not today's (the brief's requirement).

The brief's non-negotiables, built in

Regulated Australian verticals: the scorer inherits the obligations of the calls it grades.

Test data only

In place

Synthetic leads, the redacted CIMET artefact, and invented test calls. No real customer PII anywhere.

Consent is a check

In place

The recording disclaimer — including “do I have your consent” — is a critical Type A check on every call.

No card data surfaced

In place

13+ digit runs, in digits or spoken words, are masked before scoring, the LLM or disk. NBN checks the mute.

No advice, no auto-correction

In place

The app never writes to the CRM, never contacts the customer. It says what failed and where.

Rules that were live

Partial

Check version + effective date on every result. Resolving by call date is the next step.

Respect the override

In place

Reason required, machine verdict kept, append-only log, gate before → after.

Zero critical false passes

Measured against 12 hand-labelled calls. The scorer never reads the labels.

0
critical false passes

0
sales wrongly submitted

0
sales wrongly held

k 1.0
agreement with labels where the machine decided

- **Reference transcripts:** 12 calls — 11 agree, 1 escalated to QA (crosstalk), 1.2% of verdicts routed to a human.
- **Real audio (ElevenLabs Scribe):** 3 calls so far — 0 false passes, 0 wrong gates; 2 emails routed to a human by G7.
- **Caveat:** reference transcripts follow the scripts, so they're a floor; the real test is the recording on the day.

Scoring accuracy

12 hand-labelled calls · reference transcripts · Type B: deterministic-fallback · 2026-09-19T08:37:19Z

Reference transcripts

From audio (ASR)

0
critical false passes

0
sales wrongly submitted

0
sales wrongly held

100%
agreement where the machine decided

1
Cohen's kappa vs labels

1.2%
verdicts routed to a human

Unsure counts as neither agreement nor error: the system declined to guess and routed the call to QA. Labels live in data/synth/truth.json and the scorer never reads them.

Check	Critical	N	Agree	Unsure	False pass	False fail	Agree %	Kappa
account_holder	yes	12	12	0	0	0	100%	1
dmo_verbatim	yes	12	11	1	0	0	100%	1
email_captured	yes	12	12	0	0	0	100%	1
rates_and_charges	yes	12	12	0	0	0	100%	1
recording_disclaimer	yes	12	12	0	0	0	100%	1
fuel_type	-	12	12	0	0	0	100%	1
nmi	-	12	12	0	0	0	100%	1

Lead	Scenario	Label	Machine	Outcome
3613801	Clean compliant call	SUBMITTED	SUBMITTED	agree
3613802	Wrong rate quoted (27.9 vs plan 31.9)	HELD	HELD	agree
3613803	Email keyed wrong in CRM	HELD	HELD	agree
3613804	Recording disclaimer skipped	HELD	HELD	agree
3613805	DMO statement paraphrased	HELD	HELD	agree
3613806	Account holder authority never confirmed	HELD	HELD	agree
3613807	Card number read aloud (must be redacted), otherwise clean	SUBMITTED	SUBMITTED	agree

Close

Run evaluation

Proof, not claims

Everything below runs from the repo with one command each.

93

self-test assertions

Offline, ~1s: worked example, messy audio, redaction, overrides, leak guard, dialler path, G7, G8, NBN pack.

```
python run.py --selftest
```

12

labelled test calls

Wrong rate, CRM email typo, skipped disclaimer, paraphrased DMO, no authority, card read aloud, dead air, NMI mismatch, crosstalk, two criticals...

```
python tools/synth_calls.py
```

8 kHz

phone-quality audio

ElevenLabs voices or free Windows voices; 300–3400 Hz band, line noise, real overlap, real silence.

```
python tools/make_audio.py
```

E2E

through the real webhook

A dialler simulator pushes each recording, Scribe transcribes, the gate scores, the eval compares to labels.

```
python tools/dialler_sim.py --all --eval
```

Built against the five criteria

The proof is a real lead scored live — held or passed, with the exact reasoning shown.

Criterion	Weight	What we show
Scoring accuracy	30%	Eval panel: 0 critical false passes, 0 wrong gates, κ 1.0 on 12 labelled calls; per-check agreement; real Scribe audio tested.
Coverage	25%	Types A, B and C running end to end on the full Retailer 1 checklist (9 checks) and a second NBN-QA pack (12).
Traceability	20%	Quote, timestamp and check version on every result; click-to-play; one audit JSON per run; version-by-call-date next.
Gate & escalation	15%	Critical fail → HELD; unsure or < 0.60 → QA; nothing uncertain auto-passes; 5% of clean calls sampled.
Guardrails & judgment	10%	G1–G8: silence, crosstalk, mishears and screen conflicts never manufacture a false critical.

One command. No framework, no database.

Runs locally on 127.0.0.1; every external service is optional.

```
$ python run.py

CIMET QA Gate
http://127.0.0.1:8787/
Retailer 1 checklist R1-2026.09 | 9 checks
Type B: llm:claude-opus-5
ASR:     elevenlabs
Dialler webhook: POST /api/dialler/recording
```

- **Python 3.10+, standard library only** for the server, scoring engine and console.
- **Plain HTML, CSS, JavaScript** — no build step, no framework.
- **Claude (claude-opus-5)** extracts Type B values; comparisons are code.
- **ElevenLabs Scribe or Deepgram** for speech-to-text, with speaker separation.
- **No keys?** Still runs end to end on the deterministic extractor — and says so.

Setting	What it does
ANTHROPIC_API_KEY	LLM extraction for Type B
CIMET_MODEL	default claude-opus-5
ELEVENLABS_API_KEY / DEEPGRAM_API_KEY	dialler webhook (ASR)
CIMET_CONFIDENCE_FLOOR	critical floor, default 0.60
CIMET_SAMPLE_AUDIT_RATE	default 0.05
CIMET_DIALLER_TOKEN	require X-Dialler-Token
CIMET_PREScore	score new leads at startup

What's next, stated plainly

Deliberate trade-offs for a solo 12-hour build.

● Dashboards

Agent-wise daily / weekly / monthly, first-pass yield, repeat offences (3+ in 7 days). Every run's audit JSON already holds the data.

● Rules live on the call date

Versions and effective dates are stamped; resolving each check by the call date is next.

● Stereo recordings

Split agent and customer by channel for certain speaker labels (mono uses diarisation today).

● Durable job queue

The dialler queue is in memory; recordings and transcripts are already on disk.

● LLM availability

If the API fails (e.g. credit runs out), Type B falls back to the deterministic extractor and records why on every result.

● Wording drift

Synthetic calls follow the checklist wording; the real recording on the day is the better test.

The three-minute demo

A real lead scored live — held or passed, with the exact reasoning shown.

1

3613790 opens HELD

0:00

The brief's worked example: two critical fails, said vs expected on each card.

2

Click 14:02

0:30

Transcript jumps to the line; the evidence shows the quote, the plan value and the leak guard.

3

Override the rate

1:00

Reason required. Gate stays HELD until the email clears too — then SUBMITTED.

4

Switch to 3613792

1:40

Crosstalk and [inaudible]: QA, not HELD. G1 and G2 on the cards.

5

Open 3613793 (NBN)

2:10

The official artefact: \$42.90 said vs \$317 on screen, G8 → QA.

6

Accuracy

2:40

0 critical false passes, κ 1.0 — on reference and real audio.

THE ASK, ONE MORE TIME

No sale ships unscored.

Then explain exactly why, in seconds — not minutes.

HELD

QA

SUBMITTED